



# Automatic **Abuse** Detection on Social Media



# Introduction

1. Within the past year, #MeToo has gained popularity on Facebook, Reddit, and Twitter.
2. Many other hashtags have come and gone:
  - a. #YesAllWomen
  - b. #WhyIStayed
  - c. #ItsNotOkay
3. Automatic detection, classification and interpretation of personal abuse stories can help activist groups educate the public and advocate for social change in timely fashion.



## Related Work

1. **Schrading et al (2015)** assembled the Reddit Domestic Abuse Dataset. They used multiple traditional classifiers e.g., Linear SVM, logistic regression, Naive Bayes, Random Forest, etc. Highest accuracy achieved was 92.0% using Linear SVM(C=1) with N-gram features.
2. **Karlekar et al (2018)** used the same dataset with different deep learning architectures (CNN, LSTM-RNN, CNN-LSTM) They achieved an accuracy of 95.8% with CNN-LSTM model.

Model	Accuracy
Schrading et al. (2015)	92.0%
2D-CNN	92.6%
LSTM-RNN	94.5%
CNN-LSTM	<b>95.8%</b>

Table 2: Accuracy results on abuse story detection.

# Dataset

- ***redditAbuseOnlyNgrams*** This data contains a larger set of even data (1336 submissions per class), with no semantic roles or predicates. It has the variables:
  - XTrain: A list of submission title and text concatenated together, 90% training size (1202 per class).
  - XTest: A list of submission title and text concatenated together, 10% testing size (134 per class).
  - labelsTrain: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - labelsTest: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - subIdsTrain: A list of reddit submission ids, 1 entry per submission.
  - subIdsTest: A list of reddit submission ids, 1 entry per submission.

- ***redditAbuseUneven*** This data is an uneven set of data with 1336 *abuse* submissions and 17020 *non-abuse* submissions. It has the variables:
  - XTrain: A list of submission title, text, and comment data concatenated together, 85% training size.
  - XTest: A list of submission title, text, and comment data concatenated together, 15% testing size.
  - labelsTrain: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - labelsTest: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - subIdsTrain: A list of reddit submission ids, 1 entry per submission.
  - subIdsTest: A list of reddit submission ids, 1 entry per submission.

**Data Source:** Schrading et al (2015)

# Dataset

- ***redditAbuseOnlyNgrams*** This data contains a larger set of even data (1336 submissions per class), with no semantic roles or predicates. It has the variables:
  - XTrain: A list of submission title and text concatenated together, 90% training size (1202 per class).
  - XTest: A list of submission title and text concatenated together, 10% testing size (134 per class).
  - labelsTrain: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - labelsTest: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - subIdsTrain: A list of reddit submission ids, 1 entry per submission.
  - subIdsTest: A list of reddit submission ids, 1 entry per submission.



Schrading et al

- ***redditAbuseUneven*** This data is an uneven set of data with 1336 *abuse* submissions and 17020 *non-abuse* submissions. It has the variables:
  - XTrain: A list of submission title, text, and comment data concatenated together, 85% training size.
  - XTest: A list of submission title, text, and comment data concatenated together, 15% testing size.
  - labelsTrain: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - labelsTest: A list of labels (abuse or non\_abuse), 1 entry per submission.
  - subIdsTrain: A list of reddit submission ids, 1 entry per submission.
  - subIdsTest: A list of reddit submission ids, 1 entry per submission.



Karlekar at al

**Data Source:** Schrading et al (2015)

# Reference Paper

## 1. Karlekar et al (2018) : #MeToo: Neural Detection and Explanation of Language in Personal Abuse Stories

Models	
CNN	<ol style="list-style-type: none"><li>1. For each input, an embedding and a convolutional layer is applied, followed by a max-pooling layer.</li><li>2. No pre-trained word embeddings were used.</li><li>3. Filter sizes of [3, 4, 5] with 128 filters per filter were used.</li><li>4. The convolution features are then passed to a softmax layer, which outputs probabilities over two classes.</li></ol>
LSTM-RNN	<ol style="list-style-type: none"><li>1. LSTM-RNN with 128 hidden units.</li><li>2. Embedding layer followed by two LSTM hidden layers.</li><li>3. The final state is fed to a fully-connected layer and then a softmax layer, which gives the final output probabilities.</li></ol>
CNN-LSTM	<ol style="list-style-type: none"><li>1. The RNN was laid on top of CNN model</li></ol>

# Proposed Tasks

## 1. Implement the Reference Paper

**Karlekar et al (2018) : #MeToo: Neural Detection and Explanation of Language in Personal Abuse Stories**

## 2. Evaluate the model with pre-trained embeddings (*Optional*)


## 3. Use character level features in CNN (*Optional*)

# Implementation

## Libraries: Keras, Tensorflow

1. **First phase:** Pre-processing of Dataset
  - a. **Load dataset** (Data already split in training/test set)
  - b. **Visualise and Prepare Data**
2. **Second phase:** Implementation of Reference Paper
  - a. **CNN model**
  - b. **RNN - LSTM model**
  - c. **CNN - LSTM model**
3. **Third phase:** Use pre-trained word embeddings
  - a. **Glove or word2vec embeddings**
    - i. Pre-trained
    - ii. Trained on abuse dataset

Effort Estimation





# Data Pre-Processing

In [116]: *#Function to tokenize the data and add a column*

```
def tokenizer(text):  
    text = clean_text(text)  
    tokens = [word_tokenize(sent) for sent in sent_tokenize(text)]  
    tokens = list(reduce(lambda x,y: x+y, tokens))  
    tokens = list(filter(lambda token: token not in (stop + list(punctu  
    return tokens
```

x

# Data Pre-Processing

```
In [114]: #Function to remove Non-Ascii characters
#Source: https://ahmedbesbes.com/how-to-mine-newsfeed-data-and-extract-
```


```
def _removeNonAscii(s):
    return ''.join(i for i in s if ord(i)<128)
```

```
In [115]: #Function to clean text
          #Source: https://ahmedbesbes.com/how-to-mine-newsfeed-data-and-extract-
```

```
def clean_text(text):
    text = text.lower()
    text = re.sub(r"what's", "what is ", text)
    text = text.replace('ap)', '')
    text = re.sub(r"\'s", " is ", text)
    text = re.sub(r"\'ve", " have ", text)
    text = re.sub(r"can't", "cannot ", text)
    text = re.sub(r"n't", " not ", text)
    text = re.sub(r"i'm", "i am ", text)
    text = re.sub(r"\'re", " are ", text)
    text = re.sub(r"\'d", " would ", text)
    text = re.sub(r"\'ll", " will ", text)
    text = re.sub(r'\W+', ' ', text)
    text = re.sub(r'\s+', ' ', text)
    text = re.sub(r"\\", "", text)
    text = re.sub(r"\'", "", text)
    text = re.sub(r"\"", "", text)
    text = re.sub(r"^\s+", "", text)
    text = re.sub(r"^\s+|^\s*$", "", text)
    text = _removeNonAscii(text)
    text = text.strip()
    return text
```

# Data Visualisation

## 1. Data Distribution



# Data Visualisation

## 1. Most Common Words in Both Categories

### Most Common Words

```
In [25]: def keywords(label):
tokens = data[data['label'] == label]['tokens']
alltokens = []
for token_list in tokens:
    alltokens += token_list
counter = Counter(alltokens)
return counter.most_common(20)
```

```
In [26]: for label in set(data['label']):
print('label :', label)
print('top 10 keywords:', keywords(label))
print('---')
```

```
label : abuse
top 10 keywords: [('would', 3018), ('like', 2711), ('know', 2296), ('g
et', 2231), ('abuse', 1855), ('time', 1852), ('feel', 1790), ('help',
1629), ('one', 1625), ('really', 1553), ('want', 1495), ('things', 145
5), ('people', 1401), ('even', 1375), ('think', 1330), ('could', 129
0), ('never', 1181), ('life', 1170), ('years', 1131), ('going', 1129)]
---
label : non abuse
top 10 keywords: [('like', 52329), ('would', 40790), ('get', 37412),
('really', 30867), ('know', 29390), ('time', 29128), ('people', 2727
4), ('one', 27208), ('think', 22775), ('feel', 22098), ('want', 2182
0), ('go', 21760), ('good', 20596), ('going', 18402), ('something', 17
866), ('much', 17406), ('things', 17003), ('even', 16842), ('work', 16
807), ('could', 16342)]
---
```





# Data Visualisation

## 1. Finding Topics using NMF

```
In [33]: from sklearn.decomposition import NMF
```

```
In [34]: nmf = NMF(n_components=40, random_state=1, alpha=.1, l1_ratio=.5, init=  
◀────────────────────────────────────────────────────────────────────────────────▶
```

```
In [40]: feature_names = vectorizer.get_feature_names()  
no_top_words = 5  
  
for topic_idx, topic in enumerate(nmf.components_[:2]):  
    print("Topic %d:" % (topic_idx))  
    print(" | ".join([feature_names[i]  
                      for i in topic.argsort()[::-no_top_words - 1:-1]]))
```

```
Topic 0:  
like | really | know | time | think
```

```
Topic 1:  
anxiety | anxious | help | symptoms | doctor
```



# Data Visualisation

## 1. Finding Topics using NMF

In [43]: `topic_table(nmf, feature_names, no_top_words).head(5)`

Out[43]:

	topic_0:	topic_1:	topic_2:	topic_3:	topic_4:	topic_5:	topic_6:	topic_7:	topic_8:
0	like	anxiety	flair	url	job	mother	mom	anger	college
1	really	anxious	flair post	url url	work	sister	dad	angry	degree
2	know	help	add flair	picture	jobs	parents	house	rage	university
3	time	symptoms	added add	subredditlink comments	company	father	tell	control	school
4	think	doctor	bot tspq	selfie	working	brother	told	anger management	year



# Data Visualisation

## 1. Dividing topics as per abuse or non-abuse and finding intersection of topics

```
In [54]: """
showdocs(df, topics, nshow=5) is a function that gathers a number o
documents from a set of topics as a dataframe.
"""
def showdocs(df, topics, nshow=5):
    idx = df.topic == topics[0]
    for i in range(1, len(topics)):
        idx = idx | (df.topic == topics[i])
    return df[idx].groupby('topic').head(nshow).sort_values('topic')
```

```
In [55]: abuse = [0, 38, 41, 5, 43, 45, 3, 48, 29, 47]
non_abuse = [0, 1, 4, 8, 3, 16, 11, 7, 19, 18]
```

```
In [64]: showdocs(df_abuse, [37,5])['body']
```

```
Out[64]: 10      sexually abused older brothers turns dad sexua...
511     move forward wall text incoming last night mot...
582     ai child sharon jones dap kings would run rais...
618     first overnight babysitting job mother told fe...
754     abuser died sexually molested mother father so...
2       sucks man sometimes ex girlfriend used beat be...
118     deal anymore point feel thing good abused smal...
230     powerful south african psa bringing public apa...
265     abuse equal abuse equal others rant ptsd somet...
271     think need help fifteen father started asking ...
Name: body, dtype: object
```

# CNN Model

In [174]:

```
def create_conv_model():  
    model_conv = Sequential()  
    filters = 128  
    model_conv.add(Embedding(vocabulary_size, 100, input_length=6002))  
    model_conv.add(Dropout(0.2))  
    model_conv.add(Conv1D(filters, 3, activation='relu'))  
    model_conv.add(MaxPooling1D(pool_size=(3)))  
    model_conv.add(Flatten())  
    model_conv.add(Dense(1, activation='softmax'))  
    model_conv.compile(loss='binary_crossentropy', optimizer='adam',  
                       return model_conv
```

## Results

1. A model is designed that can classify stories as Abuse or Non Abuse
2. Data was visualised and techniques such as TF-IDF were employed to get important words
3. Hidden topics were decided using NMF and their relevance was shown to abuse and non-abuse stories

#ThankYou